

AUTOMATION & OPERATING MODEL BRIEF

Tender Export OS

How I turned a fragmented tender process into a governed, evidence-led operating system

Prepared for: Partner and Head of Automation, Grant Thornton

Prepared by: Raghav

Date: 30 July 2026

In one line: I have built a supervised automation system that helps a tender desk find, understand, qualify and prepare opportunities—while keeping commercially or legally binding decisions in human hands.

Why I built it

Tendering is rarely one clean process. An opportunity may begin on a portal, move into a PDF or BOQ, require supplier discovery, pass through pricing and compliance checks, and finally depend on somebody remembering a deadline or approval. In practice, this work is usually spread across browser tabs, spreadsheets, emails, folders and individual memory.

I wanted to create something broader than a tender scraper. The idea was to build the operating backbone of a tender and export desk: every opportunity receives a case ID, evidence, a current status, a clear next step and an audit trail. The system assists the team throughout the process, but it does not pretend that AI should make every decision.

What the system does

DISCOVER → VERIFY → QUALIFY → ANALYSE → SOURCE → PRICE → APPROVE → EXECUTE → LEARN

1. **Find and verify opportunities.** The Radar layer monitors tender and export sources and identifies retenders, corrigenda, repeat-buyer patterns and potentially under-served opportunities. A public listing is treated as a lead, not as proof that the opportunity is ready to pursue.
2. **Qualify quickly.** A Fast Kill stage checks the obvious constraints—eligibility, deadline, EMD, prior experience, buyer risk and product restrictions. Clear non-starters are removed early; incomplete or ambiguous cases move to a watchlist instead of being rejected on assumptions.

3. **Understand the documents.** The Deep Read stage extracts specifications, quantities, clauses, payment terms, penalties, delivery requirements and missing information from tender PDFs, BOQs, spreadsheets, corrigenda and RFQ documents.
4. **Prove that supply is viable.** The Supplier Engine uses a 5-3-2 discipline: five candidate suppliers, three different source types and two pieces of quotation proof before final pricing is considered ready.
5. **Build the commercial and compliance view.** Pricing models the complete landed-cost waterfall rather than using a single supplier number. Export cases can include EXW, FOB and CIF scenarios, while compliance remains a clearly labelled draft until an owner or qualified expert confirms it.
6. **Prepare the decision pack.** The system assembles bid packs, pricing sheets, compliance matrices, supplier comparisons, risk notes, missing-item lists and owner briefs so that the decision is made from one coherent evidence bundle.
7. **Stop for human approval.** Before any external commitment, the system creates an approval card that shows the proposed action, value, evidence, expected benefit, risks, missing information and recovery path.
8. **Track execution and outcome.** After approval, the Execution Tracker records submissions, responses, follow-ups, deadlines, receipts and final outcomes. The purpose is to learn from real results rather than from AI confidence alone.

How the system is structured

The architecture separates the work into clear layers so that research, execution, control and accountability do not become mixed together.

LAYER	ROLE IN THE OPERATING SYSTEM
Hermes control plane	Runs the operating rhythm, routes work, manages the Kanban board, surfaces blockers and approval requests, and produces the owner brief.
Specialist agents	Separate roles handle discovery, qualification, document analysis, supplier proof, pricing, compliance, pack building, approvals and execution tracking.
Python and browser capture	Perform repeatable evidence capture from known sources, document parsing, deduplication, validation and scheduled internal jobs.
Codex artifact runtime	Produces and tests the practical outputs: spreadsheets, PDFs, Word documents, dashboards, bid packs and export quotation packs.
ChatGPT research boardroom	Handles broad market research, category discovery and strategic analysis. Its output is advisory until evidence is verified locally.
Event and knowledge layer	An append-only event ledger is the canonical record. Registers, Kanban cards, reports and Drive folders are controlled views of that record rather than competing sources of truth.
Policy and approval layer	Policy-as-code, restricted tool interfaces and approval receipts keep external, financial, legal and compliance actions fail-closed.

What is genuinely different about it

- **It separates AI judgment from operational proof.** ChatGPT can discover and recommend, but deterministic tools must capture evidence before a case progresses.
- **It is a state machine, not a chain of prompts.** Every case moves through explicit statuses, gates and stop conditions.
- **Evidence has a defined maturity level.** A public listing, a downloaded document, a supplier quote and an approved action are not treated as equivalent evidence.
- **Governance is part of the architecture.** The approval boundary is enforced by policy and receipts rather than left as a reminder inside an AI prompt.
- **The system is auditable and recoverable.** The event stream can reconstruct working views and shows who or what changed a case, when it changed and what evidence supported the change.

What is automated—and what is deliberately not

Automated or automation-assisted: source monitoring, evidence capture, document extraction, deduplication, scoring, supplier research, draft pricing, draft compliance, artifact generation, readiness checks, reporting and follow-up monitoring.

Human-controlled: supplier or buyer communication, bid and quotation submission, portal uploads, payments, DSC use, final price, delivery commitments, payment terms, final classification, origin claims and legal certifications.

Current maturity

The internal operating architecture is implemented as a supervised local runtime. The repository contains the workflow definitions, event and data schemas, source adapters, parsers, policy controls, approval lifecycle, artifact generators, projection rebuilds and extensive regression and edge-case tests.

I would describe it as a working supervised automation platform—not as a finished zero-touch tender-submission bot. Credentialed portal operations and binding commercial or legal actions are intentionally outside the autonomous tool surface. That is a design choice, not an unfinished safety check.

In simple terms, I have built the digital operating backbone for a tender and export desk. It reduces dependence on scattered documents, spreadsheets and individual memory, while preserving the accountability required for high-stakes commercial work.